

L4 / Linear State Machines Formal Model

Legalese.com

November 13, 2017

Click most terms (in [this color](#)) to jump to their first underlined usage.

NTS:

- **How to read this document**
- Section and Connection
- Should I remove relievable must-next edges? **No.**

This was the idea for removing: Let e be such an edge. First, change it to a may-next edge, keeping its transition guard. So its transition guard should make its enable-times disjoint from the enable-times of other relievable must-next edges at the same state with the same role. Hmm... deriving blame (which breach state to go to) is more complicated.

Contents

1	Events, Time, Traces, Finite State Contracts	2
1.1	Execution for simple contracts	4
2	Infinite State with Global Variables	5
2.1	Execution	7
3	Event Parameters and Schema	8
3.1	Execution	9
4	May-Later and Must-Later	10

1 Events, Time, Traces, Finite State Contracts

This section defines a complete-but-limited model of contracts, called simple contracts, and also gives definitions that will be used for the full definition of contracts in Section 2.

Every contract specifies a time unit; it is the smallest unit of time that one writes constraints about or does arithmetic with. We expect it will most often be days. A time stamp is a natural number that we think of as being in units time unit.

An event E is composed of an action $action_E$, a role $role_E$, a time stamp $timestampe_E$, and optionally some parameters (but parameters will not be introduced until Section 2) $actionparams_E$. The actions and roles are fixed finite sets. In this first version of the L4 mathematical model, there is exactly one participant of each role. **All events are modelled as actions**, and a special role Env is used to model events that have no agent (i.e. role).

A trace is a sequence of events. The time stamps of the events must be nondecreasing. Thus, within the smallest unit of time, any number of events can happen; however, they are always strictly ordered. The idea here is that we want events to be strictly ordered for simplicity and to minimize the size of the space of execution traces, but if we made the time stamps strictly increasing, we would need to be working at a level of granularity for time that is at least one level smaller than the smallest unit of time that would appear in an informal version of the contract (at least when time unit = days, since contracts that use days as their minimum unit generally do not require that all events happen on different days).

A contract has a fixed finite number of states, one of which is designated the start state, and which includes at least the following:

- fulfilled
- breached(X) for each nonempty subset X of the roles. There is also an action breaches(X) for each such X , and breachEvent(X, t) is defined as the event $\langle \text{breaches}(X), \text{Env}, t \rangle$

Between any two events in a trace, the contract is in some global state G which consists of at least a state $state_G$ and a timestamp entranceTime_G (in Section 2, global variables will be added).

A **contract** has a finite directed edge-labeled multigraph¹ which we might call its **skeleton**; the nodes are the **states**, and each directed edge, which we will call an **edge**, is labeled with an **action**. The **skeleton** is the part of the **contract** that is easy to visualize. Some notation:

- For r a **role**, an **r -edge** is an **edge** whose **role** is r .
- For a an **action**, an **a -event** (**a -edge**) is an **event** (**edge**) whose **action** is a .
- For s a **state**, the **incoming s -edges** (**outgoing s -edges**) are the edges coming into (going out of) s .

Every **edge** is one of the following three types. They will be explained in more detail in the next section.

- A **may-next edge** defines permitted **events**.
- A **relievable must-next edge** defines the most used kind of obligated **events**. These are obligations that are relieved by the performance of a permitted **event** *by some other agent*.
- A **must-next edge** defines the strongest kind of obligated **events**.

Note that the events defined by **relievable must-next edges** and **must-next edges** are also considered permitted **events**.

We say that an **edge** e and an **event** $E = \langle a, r, t \rangle$ are **compatible** iff they have the same **action** a and the same **role** r . This definition will be modified in Section 2 when we add **event** parameters.

Since the environment **Env** cannot breach a contract or be *obligated* to do anything, no **Env-edge** can be a **must-next edge** or a **relievable must-next edge**. That completes the definition of the finite directed graph **skeleton** of a **contract**.

Each **edge** e is also associated with a **edge guard** $\text{edgeGuard}_e(\cdot)$ relation. For **simple contracts**, it is just a relation on **time stamps**, and an **edge** e is **enabled** upon entering a **global state** with **time stamp** t iff $\text{edgeGuard}_e(t)$ is true.²

Each **edge** e is also associated with a **deadline function** $\text{deadline}_e(\cdot)$, which yields a **deadline**. $\text{deadline}_e(t)$ is either a **time stamp** after t , or the special element ∞ . The **deadline** for an **edge** is when:

¹By this I mean there may be multiple edges from one node to another, but they must have different labels.

²Currently, LSM examples are written assuming the **edge guards** of a **state** s 's **edges** get evaluated only once upon entering the **state**. It would also be reasonable to guess that they get evaluated once per **time unit** while the **contract** is in that state. This is not ideal.

- an **enabled may-next edge** (a kind of permission) expires³.
- an **enabled must-next edge** (the strong form of obligation) causes a breach by role_e ⁴ if a **compatible event** has not been performed by the deadline.
- an **enabled relievable must-next edge** (the weak form of obligation) causes a possibly-joint breach by role_e if a **compatible event** has not been performed by the deadline **and** no other permitted **event** is performed by the deadline.

For **simple contracts**, a **deadline function** is just a function from **time stamps** to $\text{timeunit} \cup \text{timestamps}$. If d is such a function, and a state is entered at **time stamp** t , then:

- If $d(t) \in \text{timestamps}$, the deadline is $d(t)$.
- If $d(t) \in \text{timeunit}$, the deadline is $t + d(t)$.

The **edge guards** must satisfy the following conditions, which would be statically verified in a **contract-definition** language. We give the **simple contracts** definitions here, but these conditions will be used in Section 2 as well.

unambiguous absolute obligation condition: For every **time stamp** t , if some **edge guard** of a **must-next edge** evaluates to true (at t) then every other **edge guard** evaluates to false (at t).

choiceless relievable obligations condition: For every **role** r and **time stamp** t , if one of r 's **relievable must-next edges**'s **edge guards** evaluates to true (at t) then any other **relievable must-next edges** for r evaluate to false (at t).

breach or somewhere to go condition: If it is possible for all the **enabled non-Env edges** to expire simultaneously, without causing a breach (which entails that there are no enabled **must-next edges** or **relievable must-next edges**) then there must be an **Env-edge** with **deadline** ∞ .

1.1 Execution for simple contracts

A **simple contract** of course starts in its **start state**. Let E_1, E_2, \dots be a finite or infinite **trace** (recall: a sequence of **events**), as defined in Section 1. Let G_i be the **global state** that follows E_i for each i .

³Todo: expires should probably be a defined term.

⁴Which recall, in this formal model means a transition to the state $\text{breached}(\{\text{role}_e\})$

G_0 is $\langle \text{startstate}, 0 \rangle$. Let $i \geq 0$, and assume execution is defined up to entering G_i . To reduce notational clutter, let us use the aliases:

$$G = \langle s, t \rangle = G_i \quad E = \langle a, r, t' \rangle = E_i \quad G' = \langle s', t' \rangle = G_{i+1}$$

Case 1: There is some **enabled must-next edge** e in G . If there is any other **enabled edge**, then this **contract** (not just this **trace**) violates the **unambiguous absolute obligation condition**, and so is invalid.⁵

- If E is **compatible** with e and E happens within e 's deadline, then the next state must be target_e .⁶ This means E fulfilled the obligation created by e .
- Otherwise, role_e must be r and E must be $\text{breachEvent}(r, \text{deadline}_e(t) + 1)$.

Case 2: There is no **enabled must-next edge** in G . From the set of **enabled may-next edges** of s and the set of **enabled relievable must-next edges** in G , compute the deadline for each, and discard the **edges** whose deadline has passed by the time E happens;⁷ let T_p be the resulting set of **edges**. From the set of **enabled relievable must-next edges** in G , compute the deadline for each, and discard the **edges** whose deadline is not the unique minimal time stamp t^* within that set; let T_o be the resulting set, and let R be $\{\text{role}_e \mid e \in T_o\}$. Then E is either:

- An **event** compatible with some **edge** in T_p .
- $\text{breachEvent}(R, t^*)$.⁸ This means that all of the **roles** whose **enabled relievable must-next edge** expire earliest (at t^*) are jointly responsible for the breach.

The **breach or somewhere to go condition** ensures that one of those two cases will apply. In particular, it implies that at least one of T_p or R is nonempty.

2 Infinite State with Global Variables

We introduce a set of basic datatypes \mathbb{T} , which includes at least \mathbb{B} , \mathbb{N} , and \mathbb{Z} . Add to the definition of **contract** a fixed finite set of typed **global vars**. The

⁵Recall that a language (tool) for **simple contracts** will verify that such a thing can't happen.

⁶i.e. if $t' \leq \text{deadline}_e(t)$ then $s' = \text{target}_e$.

⁷i.e. discard e if $\text{deadline}_e(t) > t'$.

⁸Obviously not possible if R is empty

global vars are ordered, so we may describe their collective types as a single tuple $\text{gvartypes} \in \mathbb{T}^*$.

Add to the definition of **global state** an assignment of values to the **global vars**. We'll call such an assignment a **global vars assignment**. A particular **global vars assignment** **initvals** for the values of the **global vars** in the unique **start state** is required for a **contract**. A **contract**-definition language that introduces a notion of **contract** template should define that in terms of an assignment of parameters to **initvals**.

The **event** definition receives the following generalizations:

- Each **action** a additionally has a **global vars transform**, denoted transform_a , which is a function from $\text{gvartypes} \times \text{timestamps}$ to gvartypes .
- The definition of the **edge guard** of a a -edge is generalized: it may now depend on the values of the **global vars**; i.e. it is now a relation on $\text{timestamps} \times \text{gvartypes}$.

Now an **edgeguard** is a relation on $\text{timestamps} \times \text{gvartypes}$, and an s -edge e is **enabled** upon entering a **global state** $\langle s, t, \tau \rangle$ iff $\text{edgeGuard}_e(t, \tau)$ is true.

The three named conditions on **edge guards** are updated as follows. For every **state** s :

unique absolute obligation condition: For every **global state** G whose (local) **state** is s , if the **edge guard** of one of s 's **must-next edges** evaluates to true (on G) then every other **edge guard** of s evaluates to false.

role-unique relievable obligations condition: For every **role** r and **global state** G whose (local) **state** is s , if the **edge guard** of one of s 's **relievable must-next edges** with **role** r evaluates to true (on G) then the **edge guard** of every other of s 's **relievable must-next edges** with **role** r evaluates to false.

breach or somewhere to go condition: If it is possible for all the **enabled** non-**Env edges** to expire simultaneously, without causing a breach (which entails that there are no enabled **must-next edges** or **relievable must-next edges**) then there must be an **Env-edge** with **deadline** ∞ .

Note (probably to move to some other section or document): it will often be the case in a **contract**-definition language that we simultaneously define an **action** a and a **state** JH_a (for “ a Just Happened”, to fit its literal meaning). In this case, the incoming JH_a -edges are exactly the set of a -edges. As a convenience, a **contract**-definition language will likely allow the outgoing JH_a -edges to depend directly on a 's parameters (that is, for the **edge guard** to

depend on a 's parameters). This is merely a convenience because, as we will see when we define execution, one can achieve the same effect by introducing new **global vars** that are only used by a and JH_a ; a uses transform_a (recall, its **global vars transform**) to save its parameter values to those new **global vars**, so that the outgoing JH_a -edges can then refer to them.

2.1 Execution

Since Subsection 1.1 is short, we'll repeat essentially the entire definition of execution for **simple contracts** here, rather than say how to modify it.

Let E_1, E_2, \dots be a finite or infinite **trace** (recall: a sequence of **events**), as defined in Section 1. Let G_i be the **global state** that follows E_i for each i . A **contract** starts in its **start state**, with initial **global vars assignment** given by **initvals**.

G_0 is $\langle \text{startstate}, 0, \text{initvals} \rangle$. Let $i \geq 0$, and assume execution is defined up to entering G_i . To reduce notational clutter, let us use the aliases:

$$G = \langle s, t, \sigma \rangle = G_i \quad E = \langle a, r, t' \rangle = E_i \quad G' = \langle s', t', \sigma' \rangle = G_{i+1}$$

Case 1: There is some **enabled must-next edge** e in G . If there is any other **enabled edge**, then this **contract** (not just this **trace**) violates the **unique absolute obligation condition**, and so is invalid.⁹

- If E is **compatible** with e and E happens within e 's deadline ($t' \leq \text{deadline}_e(t)$), then the next state s' must be target_e , and σ' must be $\text{transform}_a(t, \sigma)$. This means E fulfilled the obligation created by e .
- Otherwise, role_e must be r and E must be $\text{breachEvent}(r, \text{deadline}_e(t) + 1)$ and $\sigma' = \sigma$.

Case 2: There is no **enabled must-next edge** in G . From the set of **enabled may-next edges** of s and the set of **enabled relievable must-next edges** in G , compute the deadline for each, and discard the **edges** whose deadline has passed by the time E happens¹⁰; let T_p be the resulting set of **edges**. From the set of **enabled relievable must-next edges** in G , compute the deadline for each, and discard the **edges** whose deadline is not the unique minimal time stamp t^* within that set; let T_o be the resulting set, and let R be $\{\text{role}_e \mid e \in T_o\}$. Then E is either:

⁹Recall that a language (tool) for **contracts** will verify that such a thing can't happen.

¹⁰i.e. discard e if $\text{deadline}_e(t) > t'$.

- An **event** compatible with some **edge** e in T_p . And in this case the next state s' must be target_e , and σ' must be $\text{transform}_a(t, \sigma)$
- $\text{breachEvent}(R, t^*)$.¹¹ This means that all of the **roles** whose **enabled relievable must-next edge** expire earliest (at t^*) are jointly responsible for the breach.

The **breach or somewhere to go condition** ensures that one of those two cases will apply. In particular, it implies that at least one of T_p or R is nonempty.

3 Event Parameters and Schema

Add to the definition of **contract** an assignment of types (\mathbb{T} -tuples) to the **actions**. This allows **events** to have parameters. We refer to such a type as an action-parameters domain, and the specific action-parameters domain for **action** a is paramtypes_a .

Each a -edge e gets assigned an event schema called eventschema_e . An **event schema** is a set of **events** that have the same **action**. We may think of an **event schema** as a function from $\text{gvartypes} \times \text{timestamps}$ to a set of a -events (for some fixed a). Equivalently, it is a relation on $\text{gvartypes} \times \text{timestamps} \times a\text{-events}$, and that is likely how it will be represented in a **contract**-definition language.

Non-singleton event schema are most useful for an infinite or large choice of **actions** (and, in the case of **Env-events**, for infinite or large nondeterminism).

event schema make it necessary to extend the definition of **compatible** from its previous type $\text{event} \times \text{edge}$ to $(\text{globalstate} \times \text{event}) \times \text{edge}$. We say that an **edge** e is **compatible** with $\langle G, E \rangle = \langle \langle s, t, \sigma \rangle, \langle a, r, t', \tau \rangle \rangle$ iff e is an outgoing s -edge with **action** a and **role** r , and E is in $\text{eventschema}_e(\sigma, t)$.

The three named conditions on **edge guards** are the same as before, but we add one more. We now have both **event schema** and **edge guards** as ways of constraining when an **edge** can be traversed. To reduce that redundancy we require, for every **state** s :

nonempty event schema for enabled edges: For every **global state** G whose (local) **state** is s , any **enabled** s -edge e must have eventschema_e nonempty (at G).

¹¹Obviously not possible if R is empty

3.1 Execution

Again, we elaborate the previous definition, from Subsection 2.1, of execution of a **contract** on a **trace**, but we repeat all the parts from before.

Let E_1, E_2, \dots be a finite or infinite **trace**. Let G_i be the **global state** that follows E_i for each i . A **contract** starts in its **start state**, with initial **global vars assignment** given by **initvals**.

G_0 is $\langle \text{startstate}, 0, \text{initvals} \rangle$. Let $i \geq 0$, and assume execution is defined up to entering G_i . To reduce notational clutter, let us use the following aliases, and note that we have added a forth component τ to the **event**; τ must be of type paramtypes_a .

$$G = \langle s, t, \sigma \rangle = G_i \quad E = \langle a, r, t', \tau \rangle = E_i \quad G' = \langle s', t', \sigma' \rangle = G_{i+1}$$

Case 1: There is some **enabled must-next edge** e in G . If there is any other **enabled edge**, then this **contract** (not just this **trace**) violates the **unique absolute obligation condition**, and so is invalid.¹²

- If E is **compatible** with e and E happens within e 's deadline ($t' \leq \text{deadline}_e(t)$), then the next state s' must be target_e , and σ' must be $\text{transform}_a(t, \sigma)$. This means E fulfilled the obligation created by e .
- Otherwise, role_e must be r and E must be $\text{breachEvent}(r, \text{deadline}_e(t) + 1)$ and $\sigma' = \sigma$.

Case 2: There is no **enabled must-next edge** in G . From the set of **enabled may-next edges** of s and the set of **enabled relievable must-next edges** in G , compute the deadline for each, and discard the **edges** whose deadline has passed by the time E happens¹³; let T_p be the resulting set of **edges**. From the set of **enabled relievable must-next edges** in G , compute the deadline for each, and discard the **edges** whose deadline is not the unique minimal time stamp t^* within that set; let T_o be the resulting set, and let R be $\{\text{role}_e \mid e \in T_o\}$. Then E is either:

- An **event** compatible with some **edge** e in T_p . And in this case the next state s' must be target_e , and σ' must be $\text{transform}_a(t, \sigma)$
- $\text{breachEvent}(R, t^*)$.¹⁴ This means that all of the **roles** whose **enabled relievable must-next edge** expire earliest (at t^*) are jointly responsible for the breach.

¹²Recall that a language (tool) for **contracts** will verify that such a thing can't happen.

¹³i.e. discard e if $\text{deadline}_e(t) > t'$.

¹⁴Obviously not possible if R is empty

The [breach or somewhere to go condition](#) ensures that one of those two cases will apply. In particular, it implies that at least one of T_p or R is nonempty.

4 May-Later and Must-Later

This section does not actually change the definition of [contract](#). Instead, it defines an often-useful [contract](#) structure that is likely to be supported with custom syntax in a [contract](#)-definition language.

We have so far been noncommittal about what types are available for [global vars](#). We will see later that the types strongly affect expressivity. As a special case, the reader should convince themselves that any [contract](#) that uses only boolean (or other finite domain) types can be simulated by a [simple contract](#) (using a much larger number of [states](#)).