

L4 / Linear State Machines Formal Model

Legalese.com

January 1, 2018

I recommend joining `#dsl` on our Slack and introduce yourself before delving into this. Most of the L4 documentation needs updating/improving, including this document, and Dustin or Meng will be much more motivated to prioritize that if they know there are people waiting for it.

Contents

1	Who this is for and how to read it	1
2	Events, Time, Traces, Finite State Contracts	2
2.1	Execution for simple contracts	5
3	Infinite State with Global Variables	6
3.1	Execution	8
4	Event Parameters and Schema	9
4.1	Execution	10
5	May-Later and Must-Later	11

1 Who this is for and how to read it

Please note that we have not yet taken the time to make this document as widely accessible as it will be eventually, because the contents are still changing frequently.

This document defines the programming language-independent mathematical model, (very) tentatively called “linear state machines,” that we use to define the semantics of our formal contracts language L4. It is written especially for computer scientists and mathematicians, especially Formal Verification experts, who aren’t interested in, or are turned off by, the finer details of L4. For example, consider the component of the LSM model called a “global vars transform” introduced in Section 3. It is a mathematical function, which if its domain is infinite means it’s an infinite object, which obviously won’t work for a programming language. In L4, there is an imperative-looking language for specifying such functions.

The next three sections contain complete formal contract model definitions, with Section 3 extending the model defined in Section 2, and Section 4 extending the model defined in Section 3. Section 4 is currently the most complete writeup of the L4 semantics. Our intention was to have each of those sections be independent of the others, but we’ve only succeeded partially in doing that.

Click most terms (in [this color](#)) to jump to their first underlined usage.

2 Events, Time, Traces, Finite State Contracts

This section defines a complete-but-limited model of contracts, called simple contracts, and also gives definitions that will be used for the full definition of contracts in Section 3.

Every [contract](#) specifies a time unit; it is the smallest unit of time that one writes constraints about or does arithmetic with. We expect it will most often be days or hours. A time stamp is a natural number that we think of as being in units [time unit](#).

An event E is composed of an action action_E , a role role_E , a time stamp timestamp_E , and optionally some parameters actionParams_E (but parameters will not be introduced until Section 3). The [actions](#) and [roles](#) are fixed finite sets. In this first version of the L4 mathematical model, there is exactly one participant of each [role](#).

All events in the real world are modelled as instances of [actions](#). A special [role](#) Env is used to model actions that have no subject. Each [event](#) is instantaneous (a real world event with duration is modelled by two such in-

stantaneous **events**, for the start and end), and its **time stamp** t means that the modelled real-world event happened between t and $t + 1$ **time units** since the start of the contract. There can be multiple **events** in a contract execution with **time stamp** t , and the formal model cannot assign more-precise time units to them, but the formal model does have a total order of all the **events** with timestamp t . When we need to model two real-world events as truly-simultaneous, we use one **event** to model their cooccurrence (section with example to be added below).

There can be multiple **events** in a contract execution with timestamp t , and the formal model can't assign more-precise time units to them, but the formal model does have a total order of all the **events** with timestamp t . When we need to model two truly-simultaneous informal events, we use one **event** that models their cooccurrence.

A **trace** is a sequence of **events**. The **time stamps** of the **events** must be nondecreasing. Thus, within the smallest unit of time, any number of **events** can happen; however, they are always strictly ordered. The idea here is that we want **events** to be strictly ordered for simplicity and to minimize the size of the space of execution traces, but if we made the **time stamps** strictly increasing, we would need to be working at a level of granularity for time that is at least one level smaller than the smallest unit of time that would appear in an informal version of the contract (at least when **time unit** = days, since contracts that use days as their minimum unit generally do not require that all **events** happen on different days).

A **contract** has a fixed finite number of **sections**, one of which is designated the **start section**, and which includes at least the following:

- **fulfilled**
- **breached**(X) for each nonempty subset X of the **roles**. There is also an **action** **breaches**(X) for each such X , and **breachEvent**(X, t) is defined as the **event** $\langle \text{breaches}(X), \text{Env}, t \rangle$

Between any two events in a **trace**, the **contract** is in some **global state** G which consists of at least a **section** section_G and a **time stamp** entranceTime_G (in Section 3, global variables will be added).

A **contract** has a finite directed edge-labeled multigraph¹ which we might call its **map**; the nodes are the **sections**, and each directed edge, which we

¹By this I mean there may be multiple edges from one node to another, but they must have different labels.

will call an action rule, is labeled with an **action**. The **map** is the part of the **contract** that is easy to visualize. Some notation:

- For r a **role**, an r -action rule is an action rule whose **role** is r .
- For a an **action**, an a -event (respectively a -action rule) is an **event** (respectively **action rule**) whose **action** is a .
- For s a **section**, the incoming s -action rules (outgoing s -action rules) are the action rules (edges) coming into (going out of) s .

Every **action rule** is one of the following four types. They will be explained in more detail in the next section.

- A may-next action rule defines permitted **events**.
- A relievable must-next action rule defines the most used kind of obligated **events**. These are obligations that are relieved by the performance of a permitted **event** *by some other* agent.
- A must-next action rule defines the strongest kind of obligated **events**.
- An Env action rule defines a transition that is initiated by the environment.

Note that the events defined by **relievable must-next action rules** and **must-next action rules** are also considered permitted **events**. That completes the definition of the finite directed graph “**map**” view of a **contract**.

We say that an action rule c and an **event** $E = \langle a, r, t \rangle$ are compatible iff they have the same **action** a and the same **role** r . This definition will be modified in Section 3 when we add **event** parameters.

Each **action rule** c is also associated with a relation $\text{connectionGuard}_c(\cdot)$ called its action rule guard.² For **simple contracts**, it is just a relation on **time stamps**, and an action rule c is enabled upon entering a **global state** with **time stamp** t iff $\text{connectionGuard}_c(t)$ is true.³

Each **action rule** c is also associated with a deadline function $\text{deadline}_c(\cdot)$, which yields a deadline. $\text{deadline}_c(t)$ is either a **time stamp** after t , or the special element ∞ . The **deadline** for an action rule is when:

²But note that in L4 programs, the relation may often be the trivial always-true relation.

³Currently, LSM examples are written assuming the **action rule guards** of a **section**’s **action rules** get evaluated only once upon entering the **section**. It would also be reasonable to guess that they get evaluated once per **time unit** while the **contract** is in that state. This is not ideal.

- an **enabled may-next action rule** (a kind of permission) expires⁴.
- an **enabled must-next action rule** (the strong form of obligation) causes a breach by role_c ⁵ if a **compatible event** has not been performed by the deadline.
- an **enabled relievable must-next action rule** (the weak form of obligation) causes a possibly-joint breach by role_c if a **compatible event** has not been performed by its **deadline** **and** no other permitted **event** is performed by its **deadline**.

For **simple contracts**, a **deadline function** is just a function from **time stamps** to $\text{timeunit} \cup \text{timestamps}$. If d is such a function, and a **section** is entered at **time stamp** t , then:

- If $d(t) \in \text{timestamps}$, the deadline is $d(t)$.
- If $d(t) \in \text{timeunit}$, the deadline is $t + d(t)$.

The **action rule guards** must satisfy the following conditions, which would be statically verified in a **contract-definition** language. We give the **simple contracts** definitions here, but these conditions will be used in Section 3 as well.

unambiguous absolute obligation condition: For every **time stamp** t , if some **action rule guard** of a **must-next action rule** evaluates to true (at t) then every other **action rule guard** evaluates to false (at t).

choiceless relievable obligations condition: For every **role** r and **time stamp** t , if one of r 's **relievable must-next action rules**'s **action rule guards** evaluates to true (at t) then any other **relievable must-next action rules** for r evaluate to false (at t).

breach or somewhere to go condition: If it is possible for all the **enabled non-Env action rules** to expire simultaneously, without causing a breach (which entails that there are no enabled **must-next action rules** or **relievable must-next action rules**) then there must be an **Env action rule** with **deadline** ∞ .

2.1 Execution for simple contracts

A **simple contract** of course starts in its **start section**. Let E_1, E_2, \dots be a finite or infinite **trace** (recall: a sequence of **events**), as defined in Section 2. Let G_i be the **global state** that follows E_i for each i .

⁴Todo: expires should probably be a defined term.

⁵Which recall, in this formal model means a transition to the state $\text{breached}(\{\text{role}_c\})$

G_0 is $\langle \text{startsection}, 0 \rangle$. Let $i \geq 0$, and assume execution is defined up to entering G_i . To reduce notational clutter, let us use the aliases:

$$G = \langle s, t \rangle = G_i \quad E = \langle a, r, t' \rangle = E_i \quad G' = \langle s', t' \rangle = G_{i+1}$$

Case 1: There is some **enabled must-next action rule** c in G . If there is any other **enabled action rule**, then this **contract** (not just this **trace**) violates the **unambiguous absolute obligation condition**, and so is invalid.⁶

- If E is **compatible** with c and E happens within c 's deadline, then the next state must be target_c .⁷ This means E fulfilled the obligation created by c .
- Otherwise, role_e must be r and E must be $\text{breachEvent}(r, \text{deadline}_c(t) + 1)$.

Case 2: This is rewritten case requires a modification to the **breach or somewhere to go condition**.

There are no **enabled** absolute (strong) obligations. This means that any **enabled** permitted action, **enabled** weakly-obligated action, or **enabled Env** action can occur in the next **event**. If any of those actions occur next, while the associated **deadline function** is true, then nobody breaches the contract, and there is nothing more to say. So assume otherwise.

If there is at least one permitted action that has deadline ∞ , then there also cannot be a breach, but the contract may now be stuck forever. So assume that doesn't happen either. Then at least one party breaches the contract.

Suppose that a participant p has at least one **enabled** weak obligation. For each **enabled** rule r , there is a time t_r at which all of p 's permitted actions are expired forever. If the breach happens at t_r or before, then p must be one of the breaching parties.

3 Infinite State with Global Variables

We introduce a set of basic datatypes \mathbb{T} , which includes at least \mathbb{B} , \mathbb{N} , and \mathbb{Z} . Add to the definition of **contract** a fixed finite set of typed global vars. The

⁶Recall that a language (tool) for **simple contracts** will verify that such a thing can't happen.

⁷i.e. if $t' \leq \text{deadline}_c(t)$ then $s' = \text{target}_c$.

global vars are ordered, so we may describe their collective types as a single tuple $\text{gvartypes} \in \mathbb{T}^*$.

Add to the definition of **global state** an assignment of values to the **global vars**. We'll call such an assignment a **global vars assignment**. A particular **global vars assignment** **initvals** for the values of the **global vars** in the unique **start section** is required for a **contract**; thus, our a technical definition of a **contract** is fully-instantiated, without parameters. Thus, for example, there is no **contract** representation of *the* Y-Combinator SAFE startup financing agreement, but there is a **contract** representation of every fully instantiated signed instance of it. This is not a restriction: any **contract**-definition language, such as L4, will really be a **contract**-template definition language. Making contract parameters part of the mathematical model at this point would only serve to make the model more cumbersome.⁸

The **event** definition receives the following generalizations:

- Each **action** a additionally has a **global vars transform**, denoted transform_a , which is a function from $\text{gvartypes} \times \text{timestamps}$ to gvartypes .
- The definition of the **action rule guard** of an a -action rule is generalized: it may now depend on the values of the **global vars**; i.e. it is now a relation on $\text{timestamps} \times \text{gvartypes}$.

Now a **action rule guard** is a relation on $\text{timestamps} \times \text{gvartypes}$, and an s -action rule c is **enabled** upon entering a **global state** $\langle s, t, \tau \rangle$ iff $\text{connectionGuard}_c(t, \tau)$ is true.

The three named conditions on **action rule guards** are updated as follows. For every **section** s :

unique unrelievable obligation condition: For every **global state** G whose (local) **section** is s , if the **action rule guard** of one of s 's **must-next action rules** evaluates to true (on G) then every other **action rule guard** of s evaluates to false.

role-unique relievable obligations condition: For every **role** r and **global state** G whose (local) **section** is s , if the **action rule guard** of one of s 's **relievable must-next action rules** with **role** r evaluates to true (on G) then the **action rule guard** of every other of s 's **relievable must-next action rules** with **role** r evaluates to false.

breach or somewhere to go condition: If it is possible for all the **enabled non-Env action rules** to expire simultaneously, without causing a breach (which

⁸Later, if we need to write in L^AT_EX about composing contracts, we may introduce a **contract**-template mathematical model.

entails that there are no enabled **must-next action rules** or **relievable must-next action rules**) then there must be an **Env-action rule** with **deadline** ∞ .

Note (probably to move to some other section or document): it will often be the case in a **contract-definition language** that we simultaneously define an **action** a and a **section** JH_a (for “ a Just Happened”, to fit its literal meaning). In this case, the **incoming JH_a -action rules** are exactly the set of a -**action rules**. As a convenience, a **contract-definition language** will likely allow the **outgoing JH_a -action rules** to depend directly on a ’s parameters (that is, for the **action rule guard** to depend on a ’s parameters). This is merely a convenience because, as we will see when we define execution, one can achieve the same effect by introducing new **global vars** that are only used by a and JH_a ; a uses transform_a (recall, its **global vars transform**) to save its parameter values to those new **global vars**, so that the **outgoing JH_a -action rules** can then refer to them.

3.1 Execution

Since Subsection 2.1 is short, we’ll repeat essentially the entire definition of execution for **simple contracts** here, rather than say how to modify it.

Let E_1, E_2, \dots be a finite or infinite **trace** (recall: a sequence of **events**), as defined in Section 2. Let G_i be the **global state** that follows E_i for each i . A **contract** starts in its **start section**, with initial **global vars assignment** given by **initvals**.

G_0 is $\langle \text{startsection}, 0, \text{initvals} \rangle$. Let $i \geq 0$, and assume execution is defined up to entering G_i . To reduce notational clutter, let us use the aliases:

$$G = \langle s, t, \sigma \rangle = G_i \quad E = \langle a, r, t' \rangle = E_i \quad G' = \langle s', t', \sigma' \rangle = G_{i+1}$$

Case 1: There is some **enabled must-next action rule** c in G . If there is any other **enabled action rule**, then this **contract** (not just this **trace**) violates the **unique unrelievable obligation condition**, and so is invalid.⁹

- If E is **compatible** with c and E happens within c ’s deadline ($t' \leq \text{deadline}_c(t)$), then the next state s' must be target_c , and σ' must be $\text{transform}_a(t, \sigma)$. This means E fulfilled the obligation created by c .
- Otherwise, role_c must be r and E must be $\text{breachEvent}(r, \text{deadline}_c(t) + 1)$ and $\sigma' = \sigma$.

⁹Recall that a language (tool) for **contracts** will verify that such a thing can’t happen.

Case 2: There is no **enabled must-next action rule** in G . From the set of **enabled may-next action rules** of s and the set of **enabled relievable must-next action rules** in G , compute the deadline for each, and discard the **action rules** whose deadline has passed by the time E happens¹⁰; let T_p be the resulting set of **action rules**. From the set of **enabled relievable must-next action rules** in G , compute the deadline for each, and discard the **action rules** whose deadline is not the unique minimal **time stamp** t^* within that set; let T_o be the resulting set, and let R be $\{\text{role}_c \mid c \in T_o\}$. Then E is either:

- An **event** compatible with some **action rule** e in T_p . And in this case the next state s' must be target_c , and σ' must be $\text{transform}_a(t, \sigma)$
- $\text{breachEvent}(R, t^*)$.¹¹ This means that all of the **roles** whose **enabled relievable must-next action rule** expire earliest (at t^*) are jointly responsible for the breach.

The **breach or somewhere to go condition** ensures that one of those two cases will apply. In particular, it implies that at least one of T_p or R is nonempty.

4 Event Parameters and Schema

Add to the definition of **contract** an assignment of types (\mathbb{T} -tuples) to the **actions**. This allows **events** to have parameters. We refer to such a type as an action-parameters domain, and the specific action-parameters domain for **action** a is paramtypes_a .

Each a -action rule c gets assigned an event schema called eventschema_c . An **event schema** is a set of **events** that have the same **action**. We may think of an **event schema** as a function from $\text{gvartypes} \times \text{timestamps}$ to a set of a -events (for some fixed a). Equivalently, it is a relation on $\text{gvartypes} \times \text{timestamps} \times a\text{-events}$, and that is likely how it will be represented in a **contract**-definition language.

Non-singleton event schema are most useful for an infinite or large choice of **actions** (and, in the case of **Env**-events, for infinite or large nondeterminism).

event schema make it necessary to extend the definition of **compatible** from its previous type $\text{event} \times \text{actionrule}$ to $(\text{globalstate} \times \text{event}) \times \text{actionrule}$. We say

¹⁰i.e. discard c if $\text{deadline}_c(t) > t'$.

¹¹Obviously not possible if R is empty

that an action rule c is compatible with $\langle G, E \rangle = \langle \langle s, t, \sigma \rangle, \langle a, r, t', \tau \rangle \rangle$ iff c is an outgoing s -action rule with action a and role r , and E is in $\text{eventschema}_c(\sigma, t)$.

The three named conditions on action rule guards are the same as before, but we add one more. We now have both event schema and action rule guards as ways of constraining when an action rule can be traversed. To reduce that redundancy we require, for every section s :

nonempty event schema for enabled action rules : For every global state G whose (local) section is s , any enabled s -action rule c must have eventschema_c nonempty (at G).

4.1 Execution

Again, we elaborate the previous definition, from Subsection 3.1, of execution of a contract on a trace, but we repeat all the parts from before.

Let E_1, E_2, \dots be a finite or infinite trace. Let G_i be the global state that follows E_i for each i . A contract starts in its start section, with initial global vars assignment given by initvals.

G_0 is $\langle \text{startsection}, 0, \text{initvals} \rangle$. Let $i \geq 0$, and assume execution is defined up to entering G_i . To reduce notational clutter, let us use the following aliases, and note that we have added a forth component τ to the event; τ must be of type paramtypes_a .

$$G = \langle s, t, \sigma \rangle = G_i \quad E = \langle a, r, t', \tau \rangle = E_i \quad G' = \langle s', t', \sigma' \rangle = G_{i+1}$$

Case 1: There is some enabled must-next action rule c in G . If there is any other enabled action rule, then this contract (not just this trace) violates the unique unrelievable obligation condition, and so is invalid.¹²

- If E is compatible with c and E happens within c 's deadline ($t' \leq \text{deadline}_c(t)$), then the next state s' must be target_c , and σ' must be $\text{transform}_a(t, \sigma)$. This means E fulfilled the obligation created by c .
- Otherwise, role_c must be r and E must be $\text{breachEvent}(r, \text{deadline}_c(t) + 1)$ and $\sigma' = \sigma$.

Case 2: There is no enabled must-next action rule in G . From the set of enabled may-next action rules of s and the set of enabled relievable must-next action rules in G , compute the deadline for each, and discard the action rules

¹²Recall that a language (tool) for contracts will verify that such a thing can't happen.

whose deadline has passed by the time E happens¹³; let T_p be the resulting set of **action rules**. From the set of **enabled relievable must-next action rules** in G , compute the deadline for each, and discard the **action rules** whose deadline is not the unique minimal **time stamp** t^* within that set; let T_o be the resulting set, and let R be $\{\text{role}_c \mid c \in T_o\}$. Then E is either:

- An **event** compatible with some **action rule** e in T_p . And in this case the next state s' must be target_c , and σ' must be $\text{transform}_a(t, \sigma)$
- $\text{breachEvent}(R, t^*)$.¹⁴ This means that all of the **roles** whose **enabled relievable must-next action rule** expire earliest (at t^*) are jointly responsible for the breach.

The **breach or somewhere to go condition** ensures that one of those two cases will apply. In particular, it implies that at least one of T_p or R is nonempty.

5 May-Later and Must-Later

VERY WIP

This section does not actually change the definition of **contract**. Instead, it defines an often-useful **contract** structure that is likely to be supported with custom syntax in a **contract**-definition language.

We have so far been noncommittal about what types are available for **global vars**. We will see later that the types strongly affect expressiveness. As a special case, the reader should convince themselves that any **contract** that uses only boolean (or other finite domain) types can be simulated by a **simple contract** (using a much larger number of **sections**).

¹³i.e. discard c if $\text{deadline}_c(t) > t'$.

¹⁴Obviously not possible if R is empty